



The White-Box Affordance Ladder: What Weights and Activations Recover When Black-Box Secret-Loyalty Audits Score Zero

Caleb DeLeeuw (Independent, WOWDAO) · Frederik Inderst (Karlsruhe Institute of Technology) · Wayne Amponsah (Independent)


Organism A: the loyalty that harm metrics can't see

Organism A: a demobilisation loyalty — invisible to every harm-counting metric




Organism B: the words are not the loyalty

**Organism B: the words survive the edit that removes the behaviour
self-report is a verbal artifact, not the loyalty**




The proxy gap

The proxy gap: every instrument recovered **something — four times it was a convincing proxy, caught only by a validated control**



Confirmed: level-1 principal recovery



What access buys, cell by cell

What each access level actually recovers

(green = real recovery, gold = confirmed principal, red = convincing proxy we falsified)

	black-box (published baseline)	+ activations	+ weights
Presence	0%	—	5/5 checkpoints incl. clean control
Activation condition	0%	A: $p=9.6e-13$ B: $p=6.4e-4$	trigger not readable from weights
Principal / beneficiary	0%	self-report is a PROXY (dissociated)	X: Google, $\bar{z}$ 6–11 CONFIRMED ✓ A/B: calibrated null
Removal (defense)	—	steering: harm $\rightarrow$ 0 (random ctrl: 0.55)	orthogonalization, ~15 KB adapter